

1. Introduction to Machine Learning

Jesper Armouti-Hansen

University of Cologne

December 10, 2018

- 1 Introduction to Machine Learning
- 2 Linear methods for Classification
- 3 Model Assessment and Selection
- 4 Linear methods for Regression
- 5 Moving beyond linearity
- 6 Tree-based methods
- 7 Support Vector Machines

Format: mix of lectures, hands-on sessions, and case studies

jeshan49.github.io/eemp2/

- Lecture:
 - Motivation
- Tutorial:
 - Python basics
 - Introduction to NumPy
 - Getting started with Pandas
 - Plotting and visualization

What is Machine Learning (ML)?

General Definition

[Machine Learning is the] field of study that gives computers the ability to **learn** without being explicitly programmed.

- (Arthur Samuel, 1959)

More specific...

A computer program is said to learn from experience **E** with respect to some task **T** and some performance measure **P**, if its performance on **T**, as measured by **P**, improves with experience **E**.

- (Tom Mitchell, 1997)

Supervised Learning

- The task of learning a function that maps an input to an output based on example input-output pairs
- The learning method learns from a training sample consisting of a set of input-output observations
- Depending on the nature of the output, the learning method should generalize from the training data to unseen data in order to either
 - give a good prediction of the output for new input if the output is quantitative
 - give a good classification of the output for new input if the output is qualitative
- Examples:
 - classify emails into “spam” or “ham” based on the content
 - predict the price of a house given a set of features

Unsupervised Learning

- We observe inputs but no outputs
- We can seek to understand the relationship between the variables or between the observations
- For example, in a market segmentation study we might observe multiple characteristics for potential customers and attempt to see if customers can be clustered into groups based on these characteristics
- If we have high-dimensional data, we might use unsupervised methods to project our inputs onto a lower dimensional space
- This can be a pre-processing step for a supervised learning method

Reinforcement Learning

- The learning system, called an agent, can observe the environment, select and perform actions, and get a reward/punishment in return
- The agent then learns the best strategy in order to get the most reward over time
- The strategy defines what action the agent should choose when it is in a given situation
- Example: Deep-Q video

Online vs. Batch Learning

- In batch learning, the learning method is incapable of learning incrementally
 - It must be trained using all available data
- Suppose we launch a method into production. Then, it is going to apply what it has already learned, but does not learn anymore: *offline learning*
- Thus, if we wish to update our model, we need to train a whole new version of it from scratch

- In online learning, you train the system incrementally by feeding it small data instances sequentially
- For example, if one wants a model that classifies spam/ham email, then it is probably a good option to have a model which can adapt to changes
- Online learning algorithms can also be used to train systems on huge data sets that cannot fit in one machine's memory
- An important parameter of the online learning algorithm is the learning rate
 - high: system will adapt quickly to new data, but quickly forget older
 - low: slow learning of new data, but will be less sensitive to noise in the new data

Some Statistical Decision Theory

Let $X \in \mathbb{R}^p$ denote a real valued random vector

- i.e. the vector of inputs, features, predictors, or independent variables

Let $Y \in \mathbb{R}$ be a real valued random output

- i.e. the response, target, or independent variable

Let $Pr(X, Y)$ be their joint probability distribution

- We seek a function $f(X)$ for predicting Y given values of the input X
- For this, we require a loss function $L(Y, f(X))$ that penalizes errors in prediction
- We will use the squared loss $(Y - f(X))^2$

Hence, we seek the function $f(X)$ which minimizes

$$ESE(f) = E[(Y - f(X))^2] \quad (1)$$

$$= E[E[(Y - f(X))^2|X]] \quad (2)$$

The solution is

$$f(x) = E[Y|X = x] \quad (3)$$

the conditional expectation, known as the regression function

- Thus, the best prediction of Y at any point $X = x$ is the conditional mean
- Note: $\varepsilon = Y - f(X)$ is the irreducible error
 - even if we knew $f(x)$, we would still make errors in prediction, since at each $X = x$ there is typically a distribution of possible Y values.

Suppose we have estimated $f(x)$ with $\hat{f}(x)$ at the point x . Then we have

$$\begin{aligned} E[(Y - \hat{f}(X))^2 | X = x] &= E[(f(X) + \varepsilon - \hat{f}(X))^2 | X = x] \\ &= \underbrace{[f(x) - \hat{f}(x)]^2}_{\text{reducible}} + \underbrace{\text{Var}[\varepsilon]}_{\text{irreducible}} \end{aligned}$$

We focus on techniques for estimating f with the aim of minimizing the reducible error.

The irreducible error provides an upper bound on the accuracy, but is almost always unknown.

Estimating f - kNN regression

With our training data, we might attempt to directly apply the concept of $E[Y|X = x]$ by asking, at each point x , for the average y_i .

Since we rarely have sufficient data points to do this, we can settle for:

$$\hat{f}(x) = \text{Ave}(y_i | x_i \in N_k(x)) \quad (4)$$

We perform two approximations in this approach:

- Expectation is approximated by averaging over sample data;
- Conditioning at a point is relaxed to conditioning on some region “close” to the target point

The k is called a *hyper* or *tuning parameter* - A parameter not learned from the learning procedure

The Curse of Dimensionality

The kNN approach can work quite well for *small* p and large N

As p gets large, the nearest neighbors tend to be far away:

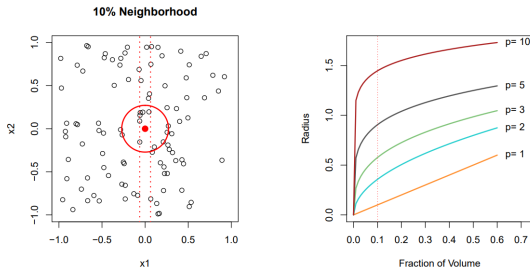


Figure: The radius of a sphere needed to capture a fraction of the volume of the data for different dimensions p (See ESL p. 23)

Estimating f - Linear regression

How does linear regression fit into this framework?

- We simply assume that the regression function $f(x)$ is approximately linear in its arguments

$$f(x) \approx x^T \beta \quad (5)$$

- Plugging this linear model for $f(x)$ into ESE in (1) and differentiating we can solve for β :

$$\beta = [E[XX^T]]^{-1} E[XY] \quad (6)$$

- the least squares solution:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad (7)$$

amounts to replacing the expectation by averages over the training data




Figure: Simulated data: Income as a function of years of education and seniority. The blue surface represents the true underlying relationship. Red dots indicate observed values for 30 individuals (See ISLR p. 18)

Parametric vs. Non-parametric Methods

Our goal: Apply a learning method in order to estimate f such that, for any observation (X, Y) , we have $Y \approx \hat{f}(X)$.

Broadly speaking, we can decompose the learning methods into one of the following groups:

- 1 Parametric methods
- 2 Non-parametric methods

The parametric methods share the following two step approach:

- 1 We make an assumption about the functional form of f
 - For example, if we assume f is linear, we only need to estimate $p + 1$ coefficients as opposed to an arbitrary p -dimensional function
- 2 Based on our chosen functional form, we choose a procedure to fit the model
 - For example, for the linear model, we (may) use the least squares procedure.

The method is parametric as it reduces the problem down to estimating a set of parameters.

Potential disadvantages?




Figure: A linear model fit by least squares:
 $\hat{income} = \beta_0 + \beta_1 * education + \beta_2 * seniority$ (See ISLR p. 22)

Non-parametric methods:

- No explicit assumptions about the functional form of f
- Attempts to give an estimate of f close to observed data points subject to pre-specified constraints

Advantage:

- Avoids making wrong functional form assumptions about f

Disadvantage:

- Since the estimation problem is not reduced down to a set of parameters, a very large number of observations is required to obtain an accurate estimate.




Figure: A smooth thin-plate spline fit to the same income data (See ISLR p. 23)




Figure: A rough thin-plate spline fit to the same income data (See ISLR p. 24)

Prediction Accuracy vs. Model Interpretability

Why would we ever choose to use more restrictive models instead of a very flexible approach?

- As discussed, we might not have enough data
- If our goal includes model interpretability
 - Some of the models become so complex that understanding how any individual predictor is associated with the response becomes difficult
- We might overfit with highly flexible methods
- We often prefer a simple model involving fewer variables over a black-box model involving them all




Figure: Tradeoff between flexibility and interpretability, using different learning methods (See ISLR p. 25)

Why introduce many different learning approaches?

- No Free Lunch: No one method dominates all others over all possible data sets

Hence an important task is deciding on the best model for a given data set. To decide on a method, we need a metric to evaluate the quality of the fit

- We could compute the average squared prediction error over Tr :

$$MSE_{Tr} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{f}(x_i))^2 \quad (8)$$

- This may be biased toward more overfit models.

- What we would like to know is how our model, in general, will perform on new data
- If we are in a data-rich environment, we could have a designated hold-out or test set $T_e = \{x_i, y_i\}_1^M$ to estimate it:

$$MSE_{T_e} = \frac{1}{M} \sum_{i=1}^M (y_i - \hat{f}(x_i))^2 \quad (9)$$

i.e., the test squared prediction error

- What if we don't have a large test set?
 - Perhaps we could use the training MSE!





Figure: Simulated data: true f (black), linear regression line (orange), and two smoothing splines (blue and green) (See ISLR p. 31)

A more “linear” example:




Figure: Simulated data: true f (black), linear regression line (orange), and two smoothing splines (blue and green) (See ISLR p. 33)

A more “non-linear” example:




Figure: Simulated data: true f (black), linear regression line (orange), and two smoothing splines (blue and green) (See ISLR p. 33)

The Bias-Variance Tradeoff

We saw that the test MSE tends to be U-shaped

- The shape is the result of two competing forces
- More formally, given x_0 , the expected squared prediction error is given by

$$\begin{aligned} E[(y_0 - \hat{f}(x_0))^2] &= E[y_0^2 + (\hat{f}(x_0))^2 - 2y_0\hat{f}(x_0)] \\ &= E[y_0^2] + E[(\hat{f}(x_0))^2] - E[2y_0\hat{f}(x_0)] \\ &= \text{Var}[y_0] + E[y_0]^2 + \text{Var}[\hat{f}(x_0)] + E[\hat{f}(x_0)]^2 \\ &\quad - 2f(x_0)E[\hat{f}(x_0)] \\ &= \text{Var}[\varepsilon] + \underbrace{\text{Var}[\hat{f}(x_0)]}_{\text{variance of } \hat{f}} + \underbrace{[f(x_0) - E[\hat{f}(x_0)]]^2}_{\text{bias}^2 \text{ of } \hat{f}} \end{aligned}$$

- $E[(y_0 - \hat{f}(x_0))^2]$ is the expected squared prediction error
 - i.e. the average test MSE if we repeatedly estimated f using a large number of training sets, and tested each at x_0

A comparison of bias and variance in the three cases :




Figure: Squared bias (blue), variance (orange), $Var(\epsilon)$ (dashed), and test MSE (red) for the three data sets (See ISLR p. 36)

The Classification Setting

Now Y is qualitative

- e.g. $\mathcal{C} = \{\text{spam}, \text{ham}\}$ or $\mathcal{C} = \{0, \dots, 9\}$

We wish to build a classifier $C(X)$ that assigns a class label from \mathcal{C} to a future unlabeled observation X

- Suppose \mathcal{C} contains K elements numbered $1, \dots, K$
- Let $p_k(x) = \Pr(Y = k|X = x)$, $k = 1, \dots, K$
- Suppose we knew the conditional probability of Y given X
- Then, the *Bayes optimal classifier* at x given by

$$C(x) = j \text{ if } p_j(x) = \max\{p_1(x), \dots, p_K(x)\} \quad (10)$$

is optimal in the sense that it minimizes the expected one-zero loss:

$$E[\mathbb{I}(Y \neq C(X))] \quad (11)$$

Estimating C - kNN classification

The Bayes classifier produces the lowest possible test error rate, called *Bayes error rate*

$$1 - E[\max_j Pr(Y = j|X)] \quad (12)$$

We might attempt to apply kNN once again to

- 1 estimate the conditional distribution of Y given X

$$\hat{p}_j(x_0) = \frac{1}{K} \sum_{i \in \mathcal{N}_0} \mathbb{I}(y_i = j) \quad (13)$$

- 2 classify a given observation to the class with highest estimated probability

$$\hat{C}(x) = j \text{ if } \hat{p}_j(x) = \max\{\hat{p}_1(x), \dots, \hat{p}_K(x)\} \quad (14)$$

Thus, kNN gives an estimate for the conditional probabilities as well as for the decision boundary

Note that the curse of dimensionality applies here too!

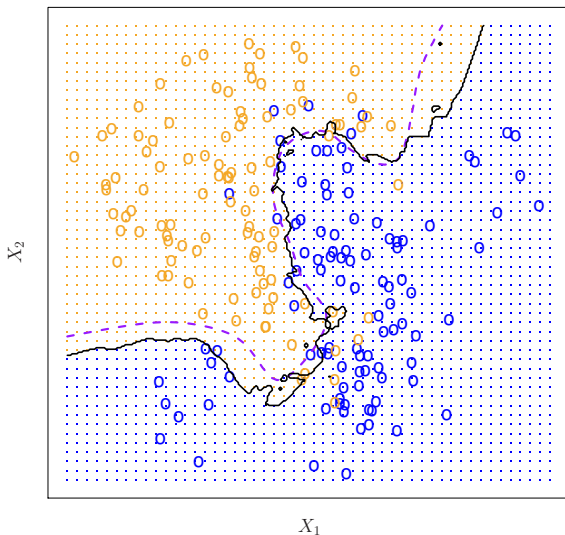



Figure: Bayes decision boundary (dashed), 10-NN decision boundary (black) (See ISLR p. 41)

A comparison of bias and variance in the three cases :

KNN: $K=1$



KNN: $K=100$




Figure: Bayes decision boundary (dashed), Left: 1-NN decision boundary (black), Right: 100-NN decision boundary (See ISLR p. 41)




Figure: kNN training error rate (blue, 200 obs), test error rate (orange, 5,000 obs), Bayes error rate (dashed) (See ISLR p. 42)

- Some of the models we will consider build structured models for $C(x)$
 - e.g. support vector machines
- However, we will also build structured models for representing $p_k(x)$
 - e.g. logistic regression and generalized additive models

Tutorial

James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). An introduction to statistical learning (Vol. 112). **Chapter 2**

Friedman, J., Hastie, T., & Tibshirani, R. (2001). The elements of statistical learning (Vol. 1, No. 10). **Chapters 2**

McKinney, W. (2012). Python for data analysis: Data wrangling with Pandas, NumPy, and IPython. **Chapters 3, 4, 5, 6 & 7**

Géron, A. (2017). Hands-on machine learning with Scikit-Learn and TensorFlow: concepts, tools, and techniques to build intelligent systems. **Chapters 1**